

Lab 001

Data cleaning and workflow [2/N]

Edward Rubin
15 January 2021

mutate

mutate

Refresher

You already know some of `mutate()`

- **Create** new variables
- **Transform** existing data into more usable forms
- **Fill in** missing values

... but it turns out **there's more!**

mutate

Pipes

But first, a reminder about pipes and the `magrittr` package.

The `%>%` operator is called a **pipe**.

```
# Equivalent:  
some_function(a, b)  
a %>% some_function(b)
```

The `%<>%` operator is an **assignment pipe**.

```
# Equivalent:  
house_df = house_df %>% mutate(n_rooms = bed + bath + other)  
house_df %<>% mutate(n_rooms = bed + bath + other)
```

mutate

'Regular' mutate

Transform a numeric variable...

```
# Convert kg to lbs: 1 kg is ~2.205 pounds  
starwars %<% mutate(weight = mass * 2.205)  
# Check results  
starwars %>% select(name, mass, weight)
```

name	mass	weight
Luke Skywalker	77	169.785
C-3PO	75	165.375
R2-D2	32	70.56

`mutate()` receives data and returns both old and new columns.

`transmute()` returns only the new columns.

mutate and family

mutate and family

Multiple columns

To `mutate()` multiple variables, you'll want one of its cousins:

- `mutate_at()` applies your function(s) to all **specified** columns.
- `mutate_if()` applies your function(s) to **"matched"** columns.
- `mutate_all()` applies your function(s) to **all** columns.

The main difference between the `mutate_*`() functions:

How do you want to select your columns?

Note: These `mutate_*`() functions will **replace** variables.

mutate and family

_at

`mutate_at()` wants three arguments:

1. `.tbl` the dataset you're trying to mutate/transform (e.g., `starwars`)
2. `.vars` the specific variables you want to transform (e.g., `mass`)
3. `.funs` the function(s) you are using for the transformation (e.g., `log`)

For example, to take the `log` of `height` and `mass`

```
# Option 1: Character vector of variable names  
starwars %>% mutate_at(.vars = c("height", "mass"), .funs = log)
```

Above: You can specify **which variables** using a character vector.

Q Where is `.tbl`?

A Implicitly defined as `starwars` via the pipe (`starwars %>%`).

mutate and family

`_at`

`mutate_at()` wants three arguments:

1. `.tbl` the dataset you're trying to mutate/transform (e.g., `starwars`)
2. `.vars` the specific variables you want to transform (e.g., `mass`)
3. `.funs` the function(s) you are using for the transformation (e.g., `log`)

For example, take the `log` of `height` and `mass`.

```
# Option 1a: Character vector of variable names
starwars %>% mutate_at(.vars = c("height", "mass"), .funs = log)
# Option 1b: Same, but don't name the arguments
starwars %>% mutate_at(c("height", "mass"), log)
```

Above: You don't have to name the arguments `.vars` and `.funs`.

mutate and family

_at

`mutate_at()` wants three arguments:

1. `.tbl` the dataset you're trying to mutate/transform (e.g., `starwars`)
2. `.vars` the specific variables you want to transform (e.g., `mass`)
3. `.funcs` the function(s) you are using for the transformation (e.g., `log`)

For example, take the `log` of `height` and `mass`.

```
# Option 2a: A "list" via the 'var' function  
starwars %>% mutate_at(.vars = vars(height, mass), .funcs = log)
```

Above: You can "list" the target variables inside `vars()`.

mutate and family

`_at`

`mutate_at()` wants three arguments:

1. `.tbl` the dataset you're trying to mutate/transform (e.g., `starwars`)
2. `.vars` the specific variables you want to transform (e.g., `mass`)
3. `.funs` the function(s) you are using for the transformation (e.g., `log`)

For example, take the `log` of `height` and `mass`.

```
# Option 2b: Variable "span" inside 'var'  
starwars %>% mutate_at(.vars = vars(height:mass), .funs = log)
```

Above: **Inside** `vars()`, define a "span" of (sequential) variables with `:`.

mutate and family

`_at`

`mutate_at()` wants three arguments:

1. `.tbl` the dataset you're trying to mutate/transform (e.g., `starwars`)
2. `.vars` the specific variables you want to transform (e.g., `mass`)
3. `.funs` the function(s) you are using for the transformation (e.g., `log`)

For example, take the `log` of `height`-related variables.

```
# Option 2c: Matching names inside 'var'.  
starwars %>% mutate_at(.vars = vars(starts_with("height")), .funs = log)
```

Above: **Inside** `vars()`, variables whose names start with `"height"`.

"Predicate" functions: `starts_with`, `ends_with`, `contains`, `matches`.

mutate and family

_at

Q How might we use `mutate_at` to help with missing values?

Ex. C-3PO and R2-D2 (among others) have `NA` for `hair_color`, but let's apply our technique for **all `"_color"` variables**.

We need two things:

1. Which **`.vars`** we will target.

```
.vars = vars(contains("_color"))
```

2. The **`function(s)`** (`.funs`) for the mutation/transformation.

```
# We'll write our own function but can rely upon 'replace_na'  
na_to_other = function(d) replace_na(data = d, replace = "other")
```

mutate and family

_at

The function `replace_na(data, replace)` comes from `tidyr`.

- `data` can be a vector or data frame
- `replace` is the replacement for those pesky `NA`s

Be careful: Just because you *can* replace `NA`s doesn't mean you *should*.

mutate and family

`_at`

Let's put it all together now.

```
# Define the function to replace NAs with "other"
na_to_other = function(d) replace_na(data = d, replace = "other")
# Use mutate_at to transform "_color" variables
starwars %>% mutate_at(
  .vars = vars(contains("_color")),
  .funs = na_to_other
)
# Check if it worked!
starwars %>% select(name, contains("_color"))
```

name	hair_color	skin_color	eye_color
Luke Skywalker	blond	fair	blue
C-3PO	other	gold	yellow
R2-D2	other	white, blue	red

mutate and family

__at

Note: We don't actually have to define our own function.

`replace_na()` just wants to know what should replace the NAs.

```
# Use mutate_at to transform "_color" variables
starwars %>% mutate_at(
  .vars = vars(contains("_color")),
  .funs = replace_na,
  replace = "other"
)
```

name	hair_color	skin_color	eye_color
Luke Skywalker	blond	fair	blue
C-3PO	other	gold	yellow
R2-D2	other	white, blue	red

mutate and family

`_if`

`mutate_if` is very similar to `mutate_at`.

- `_if` uses *logical statements* to select columns.
- *Recall:* `_at` uses *more direct statements* to select columns.

mutate and family

`_if`

`mutate_if()` wants three arguments:

1. `.tbl` the dataset you're trying to mutate/transform (e.g., `starwars`)
2. `.predicate` logical statement to select variables (e.g., `is_character`)
3. `.funs` function(s) for the transformation (e.g., `replace_na`)

For example, let's change NAs to "other" for all character variables.

```
# Replace NAs with "other" for character variables
starwars %>% mutate_if(
  .predicate = is.character,
  .funs = replace_na, replace = "other"
)
```

Above: R finds all character variables and runs `replace_na` on them.

mutate and family

`_if`

Let's change `NA`s to variables' means for `numeric` variables.

Step 1 Create a function that replaces `NA`s with the mean:

```
# Function for mean imputation/replacement  
mean_replace = function(d) replace_na(d, replace = mean(d, na.rm = T))
```

Notice: We need `na.rm = T` in `mean()` to calculate the mean when `NA`s exist.

Step 2 Use the new function `mean_replace()` with `mutate_if()`:

```
# Replace NAs with means for numeric variables  
starwars %>% mutate_if(  
  .predicate = is.numeric,  
  .funs = mean_replace  
)
```

What if we want to apply a function to **all** of our columns?

mutate and family

`_all`

`mutate_all(.tbl , .funcs)` does exactly what its name suggests.

You supply the data (`.tbl`).

`mutate_all()` transforms **all** variables in `.tbl` using `.funcs` function(s).

Now you just need a function that you want to apply to every variable...

The dplyr family

There's more!

Just as `mutate` has `mutate_at`, `mutate_if`, and `mutate_all`,

- `transmute` → `transmute_at`, `transmute_if`, and `transmute_all`
- `select` → `select_at`, `select_if`, and `select_all`
- `filter` → `filter_at`, `filter_if`, and `filter_all`
- `summarize` → `summarize_at`, `summarize_if`, and `summarize_all`
- `group_by` → `group_by_at`, `group_by_if`, and `group_by_all`

Note: `dplyr` is working to replace the "scoped verbs" (`_if`, `_at`, `_all`) with the function `across()`. But it's a bit less clean.

Finishing up

What's next?

Check out the **wine reviews**(<https://www.kaggle.com/zynicide/wine-reviews>) dataset. It has a lot of opportunities for data cleaning.

Coming soon:

- simulation, resampling, and cross validation
- `tidymodels`, `recipes` (more [here](#)), and `parsnip`
- *Related:* [tidy modeling with R](#)

Table of contents

`mutate`

- Refresher
- Pipes

`mutate_at`

- Intro
- Fundamentals
- Example

`mutate_if`

- Intro
- Fundamentals
- Example

`mutate_all`

- Intro